

# Vidit Naik

3429 Canyon Crest Dr, Riverside, CA | (951) 425-7229 | [viditnaik@gmail.com](mailto:viditnaik@gmail.com) | [linkedin.com/in/viditnaik](https://www.linkedin.com/in/viditnaik) | [viditnaik.netlify.app](https://viditnaik.netlify.app)

## Education

### University of California, Riverside

March 2025

*Master of Science in Computer Science (GPA: 3.63)*

- **Relevant Coursework:** Operating Systems, Agile Development, SDLC, Data Structures and Algorithms, Object Oriented Programming, Machine Learning, Artificial Intelligence, Networking, Design Patterns, Distributed Systems

### Vellore Institute of Technology, Chennai

August 2023

*Bachelor's in Computer Science and Engineering*

## Skills

**Languages** Java, Python, JavaScript, C++, SQL

**Database** MySQL, MongoDB, NoSQL

**Frameworks and Libraries** React.js, React Native, Hadoop, Node.js, Django, Express, RESTful APIs, Flask, SpringBoot

**Testing** PyTest, JUnit, Selenium

**AI and ML** Large Language Models (Qwen, BERT, Llama3), RAG, PyTorch, Fine-Tuning (LoRA), Ollama

**Cloud and DevOps** AWS (S3, Glue, Lambda, EC2), Docker, Kubernetes, CI/CD Pipelines, Jenkins, Git/GitHub, Jira

**Tools and Methodologies** Microservices, Tableau, Distributed Systems, RESTful API, Agile/Scrum, Jira, OAuth 2.0, Linux, Unix Shell Scripts

## Experience

### Shifa Precision Inc.

July 2025 - Present

*Software Engineer - AI/ML*

*Boston, MA*

- Leading the backend development of Project Oasis - a platform leveraging AWS and Neo4j to generate real time patient “digital twins”.
- Engineering a scalable data pipeline to unify biomedical data (ClinVar, PubMed, openFDA) into a knowledge graph with over 1M+ nodes and 10M+ relationships
- Implemented backend services integrating LLM APIs for biomedical data extraction and supported deployment of a GNN-based prediction engine

### Center for Robotics and Intelligent Systems

October 2024 - March 2025

*Student Researcher*

*Riverside, CA*

- Developed an AI-powered system that integrates drones with Large Language Models (LLMs) for user-driven control through natural language.
- Utilized Retrieval-Augmented Generation (RAG) to provide the LLM with contextual knowledge from drone manuals, increasing command accuracy by 68%.
- Collaborated in a Scrum-based Agile development process to design features with a core commitment to quality and scalability.

### StuDetails

July 2020 - February 2021

*Software Engineer*

*Noida, India*

- Architected and developed a full-stack web application using React and Flask, enabling 50+ enterprise clients to manage scheduling and inspections, resulting in a 20% increase in operational efficiency, supporting over 5,000 monthly active users.
- Led cross-functional collaboration to optimize AWS S3 data workflows using AWS Glue, improving data accuracy by 25% and enabling real-time integration with client-facing systems.
- Implemented automated testing and CI/CD pipelines, cutting deployment failures by 30% for apps processing 1TB+ data monthly.

## Projects

### InsureSearch: RAG based AI Chatbot

- Developed and deployed a full-stack AI chatbot using a React frontend, a Node.js/Express.js backend, and a MongoDB database.
- Optimized the backend to reduce token usage by 65% through a top-k BERT search implementation, significantly improving response times.
- Enhanced LLaMA-3 8B with 4-bit quantization and LoRA adapters, fine-tuning on insurance policy data to boost accuracy for user prompts. Handled the source control of the project through Git.

### CitySafe: Chicago Crime Insights Dashboard

- Built an analytics dashboard processing 1.5M+ records using SQL, Spark, and PostGIS with ETL pipelines for automated data extraction and crime trend insights.
- Designed geospatial visualizations with Python libraries and React.js, enabling self-service monitoring of crime hotspots.
- Implemented a Python-based RESTful API with optimized data models and indexing, which reduced query response time by 25% and efficiently served data to a React.js frontend.

### CreateFlow - CalHacks 11.0

- Designed a multi-agent system with LangGraph linking AI agents for content, scheduling, & analytics, cutting creation time by 50%
- Enhanced LLaMA-3 8B with 4-bit quantization and LoRA adapters, fine-tuning on LinkedIn data to boost engagement predictions